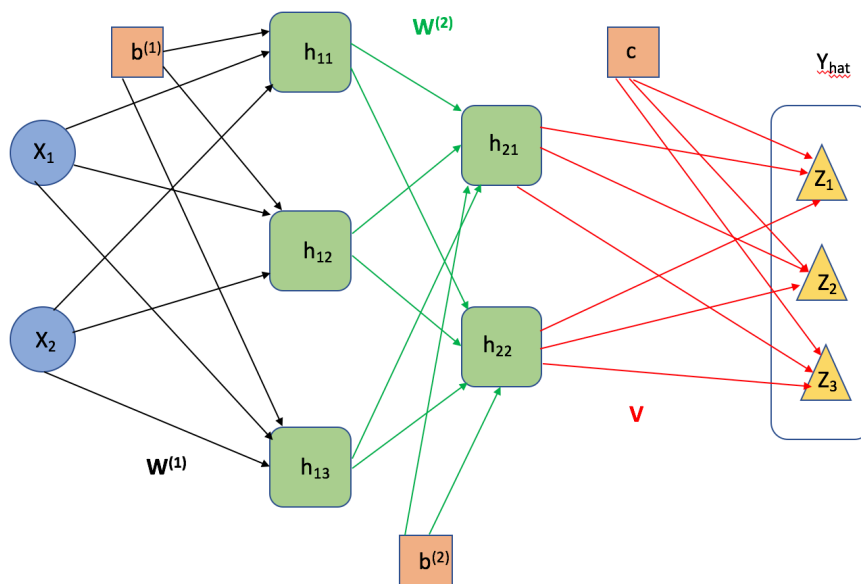


Brody Vogel Homework # 1

October 12, 2018

```
In [332]: import numpy as np
import matplotlib.pyplot as plt
```



1.1

$$\begin{aligned}h_{11} &= \max(0, \mathbf{W}^{(1)}_{11}X_1 + \mathbf{W}^{(1)}_{21}X_2 + b^{(1)}_1) \\h_{12} &= \max(0, \mathbf{W}^{(1)}_{12}X_1 + \mathbf{W}^{(1)}_{22}X_2 + b^{(1)}_2) \\h_{13} &= \max(0, \mathbf{W}^{(1)}_{13}X_1 + \mathbf{W}^{(1)}_{23}X_2 + b^{(1)}_3)\end{aligned}$$

$$\begin{aligned}h_{21} &= \max(0, \mathbf{W}^{(2)}_{11}h_{11} + \mathbf{W}^{(2)}_{21}h_{12} + \mathbf{W}^{(2)}_{31}h_{13} + b^{(2)}_1) \\h_{22} &= \max(0, \mathbf{W}^{(2)}_{12}h_{11} + \mathbf{W}^{(2)}_{22}h_{12} + \mathbf{W}^{(2)}_{32}h_{13} + b^{(2)}_2)\end{aligned}$$

$$\begin{aligned}Z_1 &= \max(0, \mathbf{V}_{11}h_{21} + \mathbf{V}_{21}h_{22} + c_1) \\Z_2 &= \max(0, \mathbf{V}_{12}h_{21} + \mathbf{V}_{22}h_{22} + c_2) \\Z_3 &= \max(0, \mathbf{V}_{13}h_{21} + \mathbf{V}_{23}h_{22} + c_3)\end{aligned}$$

$$\mathbf{Y}_{\text{hat}} = [(e^{Z_1}) / (\sum_k(e^{Z_k})), (e^{Z_2}) / (\sum_k(e^{Z_k})), (e^{Z_3}) / (\sum_k(e^{Z_k}))]$$

Or, with matrix algebra:

$$\begin{aligned}\mathbf{h}_1 &= \max(0, \mathbf{XW}^{(1)} + \mathbf{b}^{(1)}) \\ \mathbf{h}_2 &= \max(0, \mathbf{h}_1\mathbf{W}^{(2)} + \mathbf{b}^{(2)}) \\ \mathbf{Z} &= \max(0, \mathbf{h}_2\mathbf{V} + \mathbf{c}) \\ \mathbf{Y}_{\text{hat}} &= (\mathbf{e}^{\mathbf{Z}}) / (\sum(\mathbf{e}^{\mathbf{Z}}))\end{aligned}$$

1.2

In [333]: # 1.3

```
# ReLU activation function
def ReLu(vector):
    return(np.maximum(0, vector))

# softmax output function
def softmax(vector):
    # took some of this from StackOverflow (didn't know about 'np.exp')
    e_x = np.exp(vector)
    return e_x / e_x.sum(axis = 0)

# the neural net
def ff_nn_2_ReLu(input_vector, weights1, weights2, weights3, bias1, bias2, bias3):
    # compute the output step-by-step
    # h1 = W1X + b1^T (3 x 3)
    h1 = ReLu(weights1.dot(input_vector) + bias1)
    # h2 = W2h1 + b2^T (2 x 3)
    h2 = ReLu(weights2.dot(h1) + bias2)
    # Z = Vh2 + c^T (3 x 3)
    z = weights3.dot(h2) + bias3
    # y_hat
    y_hat = softmax(z)

    return(y_hat)
```

In [334]: # 1.4

```

# compute the specific output for the supplied input
X = np.array([[1, 0, 0],
               [-1, -1, 1]])

W1 = np.array([[1, 0],
               [-1, 0],
               [0, .5]])

W2 = np.array([[1, 0, 0],
               [-1, -1, 0]])

V = np.array([[1, 1],
               [0, 0],
               [-1, -1]])

b1 = np.array([0, 0, 1]).reshape(3, 1)
b2 = np.array([1, -1]).reshape(2, 1)
c = np.array([1, 0, 0]).reshape(3, 1)

ff_nn_2_ReLu(X, W1, W2, V, b1, b2, c)

Out [334]: array([[ 0.94649912,  0.84379473,  0.84379473],
                  [ 0.04712342,  0.1141952 ,  0.1141952 ],
                  [ 0.00637746,  0.04201007,  0.04201007]])

```

2.1

$$\frac{\partial}{\partial x}(1-x)^2 + 100(y-x^2)^2 = 2(1-x)(-1) + 200(y-x^2)(2x) = -2 + 2x - 400(xy - x^3)$$


$$\frac{\partial}{\partial y}(1-x)^2 + 100(y-x^2)^2 = 200(y-x^2)(1) = 200(y-x^2)$$

In [335]: # 2.2

```

# adapted from the Gradient_Descent notebook
# contours are hard to capture, so we have to zoom way in on delta
delta = .01
x = np.arange(-1, 1, delta)
y = np.arange(-1, 1, delta)
X, Y = np.meshgrid(x, y)
# here's the Rosenbrock Function
Z = (1-X)**2 + 100*(Y-X**2)**2
fig, ax = plt.subplots()
CS = ax.contour(X, Y, Z)

```



In [336]: # 2.3

```
# function for computing the instantaneous gradient
# also adapted from Gradient_Descent notebook
def grad_f(vector):
    x, y = vector
    # partial derivative with respect to x
    df_dx = -2 + 2*x - 400*(x*y-x**3)
    # partial derivative with respect to y
    df_dy = 200*(y - x**2)
    # gradient of the two partial derivatives
    return np.array([df_dx, df_dy])

# function for the gradient descent algorithm
# again, adapted from the Gradient_Descent notebook
def grad_descent(starting_point=None, iterations=20, learning_rate=12):
    if starting_point:
        point = starting_point
    else:
        # have to start in a small interval to keep things in check
        point = np.random.uniform(-1,1,size=2)
    trajectory = [point]

    for i in range(iterations):
        grad = grad_f(point)
        point = point - learning_rate * grad
```

```

        trajectory.append(point)
    return np.array(trajectory)

np.random.seed(10)
# learning rate of .005
traj = grad_descent(iterations=200, learning_rate = .005)
traj1 = grad_descent(iterations=200, learning_rate = .0005)
traj2 = grad_descent(iterations=200, learning_rate = .008)


trajs = [traj, traj1, traj2]
lrates = [.005, .0005, .008]

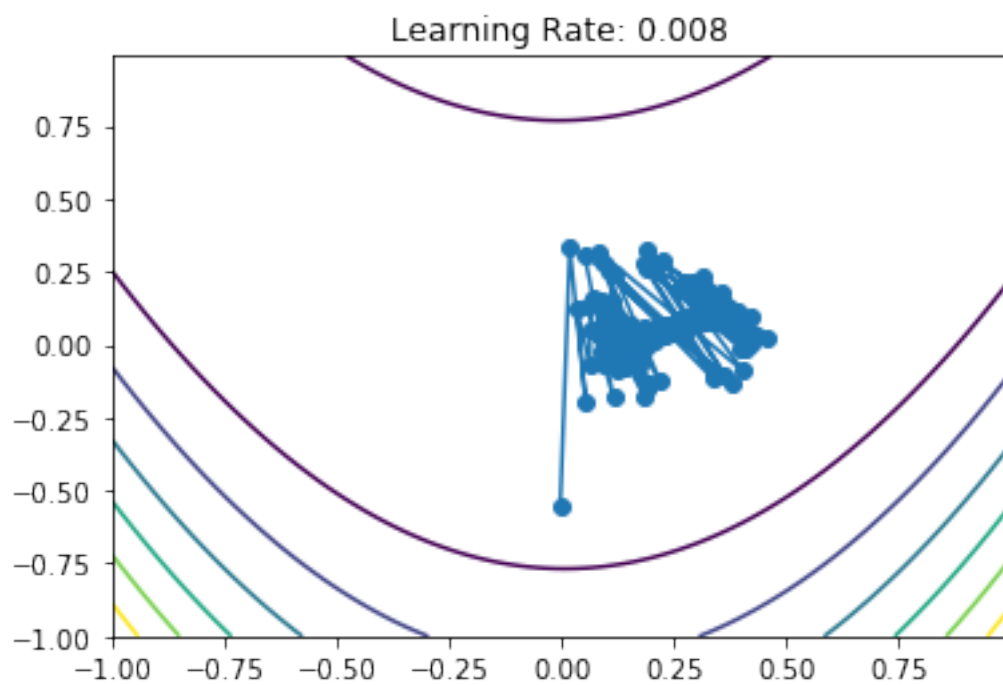

for num in range(3):

    t = trajs[num]
    lr = lrates[num]
    fig, ax = plt.subplots()
    CS = ax.contour(X, Y, Z)
    x= t[:,0]
    y= t[:,1]
    plt.title("Learning Rate: " + str(lr))
    plt.plot(x,y,'-o')

# Each of the learning rates - with 200 iterations - end up in about the
# same place, albeit after much different paths.
# The largest tested learning rate (.008), especially, showed erratic
# behavior, which makes me think it overshoots the minimum in the beginning.

```





In [337]: # 2.4

```
# gradient descent with momentum algorithm
# adapted from the Gradient_Descent notebook
def grad_descent_with_momentum(starting_point=None, iterations=10, alpha=.9, epsilon=.0001):
    if starting_point:
        point = starting_point
    else:
        point = np.random.uniform(-1,1,size=2)
    trajectory = [point]
    # initialize the velocity vector
    v = np.zeros(point.size)

    for i in range(iterations):
        # call the same instantaneous-gradient-finding function
        grad = grad_f(point)
        # create the update to the point in the gradient descent with momentum
        # algorithm by summing the instantaneous gradient and the past gradients
        # in the velocity vector
        v = alpha*v + epsilon*grad
        # update the point in the gradient descent trajectory
        point = point - v
        trajectory.append(point)
    return np.array(trajectory)

# test different values for epsilon and alpha

np.random.seed(10)
traj = grad_descent_with_momentum(iterations=50, epsilon=.0002, alpha=.005)
traj1 = grad_descent_with_momentum(iterations=50, epsilon=.0002, alpha=.002)
traj2 = grad_descent_with_momentum(iterations=50, epsilon=.0002, alpha=.001)
traj3 = grad_descent_with_momentum(iterations=50, epsilon=.0002, alpha=.005)
traj4 = grad_descent_with_momentum(iterations=50, epsilon=.0005, alpha=.005)
traj5 = grad_descent_with_momentum(iterations=50, epsilon=.001, alpha=.005)

trajs = [traj, traj1, traj2, traj3, traj4, traj5]
eps = [.0002, .0002, .0002, .0002, .0005, .001]
alphs = [.005, .002, .001, .005, .005, .005]

for num in range(6):
    t = trajs[num]
    e = eps[num]
    a = alphs[num]

    fig, ax = plt.subplots()
    CS = ax.contour(X, Y, Z)
    x= t[:,0]
```


```



y= t[:,1]
plt.title("Epsilon: " + str(e) + "; Alpha: " + str(a))
plt.plot(x,y,'-o')


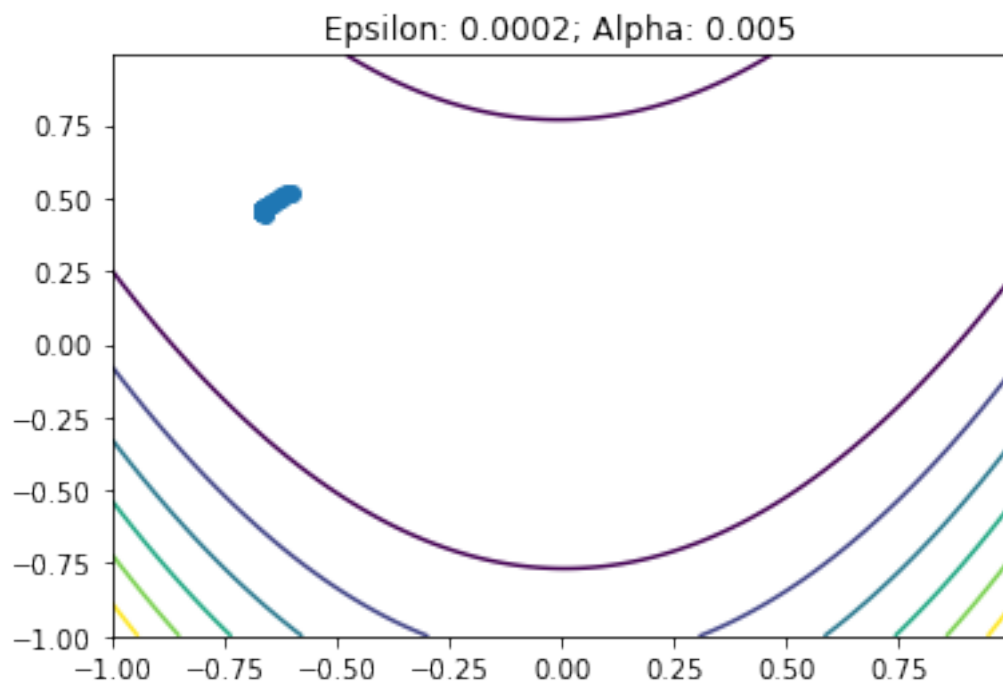
```


*# Each of the different tested hyperparameters for tuning the weight of the
current gradient (alpha) end up sending the trajectory in the same
direction, but the larger learning rates look like they make it head there
faster.*

*# The same can be said for the tested epsilon parameters that control
the weight given to the velocity vector. It looks like they're all
headed towards a minimum, but the ones that use larger parameters look
like they'd get there faster.*









3.1

Had to rearrange things a bit from Nielsen's notation to account for the data.

$$\frac{\partial L}{\partial c} = \sum (\hat{y} - y)$$

$$\frac{\partial L}{\partial V} = (\hat{y} - y) \cdot out_2^T$$

$$\frac{\partial L}{\partial b_2} = \sum (V^T \cdot (\hat{y} - y)) \circ out_2$$

$$\frac{\partial L}{\partial W_2} = (V^T \cdot (\hat{y} - y)) \circ out_2 \cdot out_1^T$$

$$\frac{\partial L}{\partial b_1} = \sum W_2^T \cdot (V^T \cdot (\hat{y} - y)) \circ out_2 \circ out_1$$

$$\frac{\partial L}{\partial W_1} = (W_2^T \cdot (V^T \cdot (\hat{y} - y)) \circ out_2 \circ out_1) \cdot X^T$$

$$out_1 = \text{ReLU}(W_1 \cdot X + b_1)$$

$$out_2 = \text{ReLU}(W_2 \cdot out_1 + b_2)$$

In [338]: # 3.2

```
def grad_f(param_vec, x, y):

    # unpack the parameters
    W1 = param_vec[0:6].reshape(3, 2)
    W2 = param_vec[6:12].reshape(2, 3)
    v = param_vec[12:18].reshape(3, 2)
    b1 = param_vec[18:21].reshape(3, 1)
    b2 = param_vec[21:23].reshape(2, 1)
    c = param_vec[23:26].reshape(3, 1)

    # forward pass
    a1 = W1.dot(x) + b1
    H1 = ReLu(a1)
```

```

a2 = W2.dot(H1) + b2
H2 = ReLu(a2)
Z = v.dot(H2) + c
y_hat = softmax(Z)

# the partials, as described above
d_v = (y_hat - y).dot(H2.T) # d_V
d_c = (y_hat - y).sum(axis = 1) # d_c
d_W2 = ((V.T.dot((y_hat - y))) * (H2 > 0)).dot(H1.T) # d_W2
d_b2 = ((V.T.dot((y_hat - y))) * (H2 > 0)).sum(axis = 1) # d_b2
d_W1 = (W2.T.dot((V.T.dot((y_hat - y))) * (H2 > 0)) * (H1 > 0)).dot(x.T) # d_W1
d_b1 = (W2.T.dot((V.T.dot((y_hat - y))) * (H2 > 0)) * (H1 > 0)).sum(axis = 1) # d_b1

# repack the parameters
z = [d_W1, d_W2, d_b1, d_b2, d_v, d_c]
z = [param.flatten() for param in z]
z = np.concatenate(z)

# return the gradient
return(z)

def grad_descent(x, y, iterations=10, learning_rate=1e-2):
    # initialize the weights and biases
    v = np.random.uniform(-.1, .1, size = 6).reshape(3, 2)
    w2 = np.random.uniform(-.1, .1, size = 6).reshape(2, 3)
    w1 = np.random.uniform(-.1, .1, size = 6).reshape(3, 2)
    c = np.random.uniform(-.1, .1, size = 3).reshape(3, 1)
    b2 = np.random.uniform(-.1, .1, size = 2).reshape(2, 1)
    b1 = np.random.uniform(-.1, .1, size = 3).reshape(3, 1)

    # pack the parameters
    point = [param.flatten() for param in [w1, w2, v, b1, b2, c]]
    point = np.concatenate(point)

    # initiate the trajectory and loss vectors
    trajectory = [point]
    losses = [loss(y, y_hat(x, point))]

    # perform the gradient descent algorithm
    for i in range(iterations):
        grad = grad_f(point, x, y)
        point = point - learning_rate * grad
        trajectory.append(point)
        losses.append(loss(y, y_hat(x, point)))
    return (np.array(trajectory), losses)

# categorical cross-entropy loss function

```

```

def loss(y, y_hat):
    # cross entropy
    tot = y * np.log(y_hat)
    return -tot.sum()

# function for computing a forward pass
def y_hat(input_vector, param_vec):
    # unpack the parameters
    weights1 = param_vec[0:6].reshape(3, 2)
    weights2 = param_vec[6:12].reshape(2, 3)
    weights3 = param_vec[12:18].reshape(3, 2)
    bias1 = param_vec[18:21].reshape(3, 1)
    bias2 = param_vec[21:23].reshape(2, 1)
    bias3 = param_vec[23:26].reshape(3, 1)

    # compute the output step-by-step
    #  $h_1 = W_1X + b_1^T$  (3 x 3)
    h1 = ReLu(weights1.dot(input_vector) + bias1)
    #  $h_2 = W_2h_1 + b_2^T$  (2 x 3)
    h2 = ReLu(weights2.dot(h1) + bias2)
    #  $Z = Vh_2 + c^T$  (3 x 3)
    z = weights3.dot(h2) + bias3
    # y_hat
    y_hat = softmax(z)
    # return the estimate
    return(y_hat)

```

In [339]: # 3.3

function for generating the Gaussian data, taken from the example notebook

```

import pandas as pd
def gen_gmm_data(n = 999, plot=False):
    # Fixing seed for repeatability
    np.random.seed(123)

    # Parameters of a normal distribuion
    mean_1 = [0, 2] ; mean_2 = [2, -2] ; mean_3 = [-2, -2]
    mean = [mean_1, mean_2, mean_3] ; cov = [[1, 0], [0, 1]]

    # Setting up the class probabilities
    n_samples = n
    pr_class_1 = pr_class_2 = pr_class_3 = 1/3.0
    n_class = (n_samples * np.array([pr_class_1, pr_class_2, pr_class_3])).astype(int)

    # Generate sample data
    for i in range(3):
        x1,x2 = np.random.multivariate_normal(mean[i], cov, n_class[i]).T
        if (i==0):

```

```

        xs = np.array([x1,x2])
        cl = np.array([n_class[i]*[i]])
    else:
        xs_new = np.array([x1,x2])
        cl_new = np.array([n_class[i]*[i]])
        xs = np.concatenate((xs, xs_new), axis = 1)
        cl = np.concatenate((cl, cl_new), axis = 1)

    # One hot encoding classes
    y = pd.Series(cl[0].tolist())
    y = pd.get_dummies(y).as_matrix()

    # Normalizing data (prevents overflow errors)
    mu = xs.mean(axis = 1)
    std = xs.std(axis = 1)
    xs = (xs.T - mu) / std

    return xs, y, cl

```

```


In [342]: x,y,cl = gen_gmm_data(500)
plt.scatter(x[:,0], x[:,1], c=cl)

```

```

# I transposed x and y to make the computations easier, since the dimensions in the
# assignment are a little different from those in the example.
x = x.T
y = y.T


```



In [354]: # 3.4

```
traj, losses = grad_descent(x, y, iterations=400, learning_rate=.0002)
# here's the plot of the losses along the parameters trajectory - looks like it's lea
# and I had to play with the learning rate a lot
plt.plot(losses)
```

Out [354]: [matplotlib.lines.Line2D at 0x119eb0550>]



In [356]: # 3.5

```
# another gradient descent with momentum function
def grad_descent_with_momentum_2(x, y, iterations=10, alpha=.9, epsilon=10):
    # initialize the weights and biases
    v = np.random.uniform(-.1, .1, size = 6).reshape(3, 2)
    w2 = np.random.uniform(-.1, .1, size = 6).reshape(2, 3)
    w1 = np.random.uniform(-.1, .1, size = 6).reshape(3, 2)
    c = np.random.uniform(-.1, .1, size = 3).reshape(3, 1)
    b2 = np.random.uniform(-.1, .1, size = 2).reshape(2, 1)
    b1 = np.random.uniform(-.1, .1, size = 3).reshape(3, 1)

    # pack the parameters
    point = [param.flatten() for param in [w1, w2, v, b1, b2, c]]
    point = np.concatenate(point)
    # initialize the velocity vector
    v = np.zeros(point.size)
```

```

# initiate the trajectory and loss vectors
trajectory = [point]
losses = [loss(y, y_hat(x, point))]

# perform the gradient descent algorithm
for i in range(iterations):
    grad = grad_f(point, x, y)
    # create the update to the point in the gradient descent with momentum
    # algorithm by summing the instantaneous gradient and the past gradients
    # in the velocity vector
    v = alpha*v + epsilon*grad
    # update the point in the gradient descent trajectory
    point = point - v
    trajectory.append(point)
    losses.append(loss(y, y_hat(x, point)))
return (np.array(trajectory), losses)

```

```

In [359]: # It looks like, with momentum, it learns about the same, and
# maybe a little bit less efficiently.
# That is, it took longer to start descending down the loss
# slope, and it didn't traverse
# down the slope as quickly. Although, again,
# this took a lot of parameter "tuning" (guessing)
traj1, losses1 = grad_descent_with_momentum_2(x, y, iterations = 400, alpha = .001, epsilon = .01)
plt.plot(losses1)

```

```

Out [359]: [<matplotlib.lines.Line2D at 0x11b19d9b0>]

```

